

Secure Cloud Computing: Encrypted Databases and Searchable Symmetric Encryption

Riccardo Gennaro

*Faculty of Electrical Engineering, Mathematics and Computer Science
University of Twente*

Péter Szelecz

*Faculty of Electrical Engineering, Mathematics and Computer Science
University of Twente*

1. Proving Information Leakage of an SSE Scheme

1.1 Approach for Proving Information Leakage in SSE Schemes

As described in the assignment text, in order to prove an upper bound for the information leaked, or to quantifying the information leakage for an SSE scheme, we base our approach on simulation-based experiments as per Curtmola et al. [article:curtmola]. In particular, we base our approach on the simulation for proving non-adaptive semantic security.

We perform the following:

- **Define the Leakage Function L :** the leakage function expresses which information are exposed to an adversary during the execution of a protocol based on the given SSE schema. In particular, this function specifies the nature of the leaked information, e.g. number of document, access frequencies, etc. This leakage is then given as input to the simulator.
- **Construct a simulator S for the protocol with the stated leakage L :** the simulator is an algorithm that simulates the adversary view of the real protocol. This algorithm takes as input the leakage of the SSE schema and simulate an output that needs to appear indistinguishable from the output of the execution of the real protocol.
- **Prove indistinguishability of the execution from the adversary perspective:** the adversary (or distinguisher) send the challenger a set of query. The challenger execute either the real or the simulated protocol. Indistinguishability is proven if the adversary is not able to distinguish between the real and the simulated output with a negligible advantage ϵ .

1.2 Leakage statement

The provided SSE scheme present the following elements in the leakage:

- **Number of documents $|\mathcal{D}|$:** since by construction documents D_i are mapped one-to-one to encrypted keywords sets C_i , the number of C_i sets is equal to the number of documents. Since the index is $\mathcal{I} = C_1, \dots, C_d$, where d is the number of documents, the number of documents is leaked.
- **Number of keywords per document $n_i = |C_i|$:** since the index \mathcal{I} contains the set of encrypted keywords C_i , the number of keywords associated to a document C_i is leaked.
- **Keyword equality - number of distinct keywords:** since we the keywords are encrypted using a deterministic pseudorandom function $F : 0, 1^\lambda \times 0, 1^* \rightarrow 0, 1^n$, for equal keywords $p_i = p_j$ the pseudorandom function F will return equal results $\gamma_i = F(k, p_i) = F(k, p_j) = \gamma_j$. This leaks the number of distinct keywords.

- **Search pattern** $SP(\mathbf{q})$: the adversary is able to determine what queries are for the same keyword since the search tokens τ_q are deterministically computed using pseudorandom permutation F .
- **Access pattern** $AP(\mathbf{q})$: the adversary is able to determine what document identifiers are returned for a given query q_i as query result $D(q_i)$.

■ 1.3 Proof sketch

Given the access pattern $AP(\mathbf{q})$, for all sets $D(q_i) \in AP(\mathbf{q})$ with i s.t. $1 \leq i \leq q$, for every document identifier $j \in D(q_i)$, insert in the set of encrypted keywords C'_j of document D_j the simulated encrypted keyword γ'_j sampled uniformly and randomly as $\gamma'_j \xleftarrow{\$} \{0,1\}^n$, instead of being computed with pseudorandom permutation F . As a result, we obtain simulated index $\mathcal{I}' = C'_1, \dots, C'_d$,

For every $q_j \in SP(\mathbf{q})$, set search token $\tau'_j = \gamma'_j$ instead of computing it with pseudorandom permutation F .

We compare the simulated parameters with the real ones:

- **\mathcal{I} and \mathcal{I}'** : \mathcal{I} is the collection C'_1, \dots, C'_d containing the sets of the keywords of every document encrypted with deterministic pseudorandom permutation F . If collection $AP(\mathbf{q})$ is non-empty, \mathcal{I}' will contain the collection C'_1, \dots, C'_d of the sets of the simulated keywords of the documents indexed in $AP(\mathbf{q})$. The simulated encrypted keywords are sampled uniformly and randomly and with length n , matching the length of the real keywords. Since the key used with F is not known by the distinguisher, it is not able to distinguish C_i from C'_i .
- **τ_j and τ'_j** : τ_j is computed via the deterministic F function and its value will match one of the real encrypted keywords. The value of τ'_j is assigned in such a way that for query q_j in $SP(\mathbf{q})$ it will be equal to simulated encrypted keyword γ'_j .

Since the structure is preserved, the only way for a distinguisher to distinguish between real and simulated parameters is to be able to distinguish between the output of the random and uniform sampling and the output of the PRP function that is considered secure.

2. Attacking Property-preserving Encryption

2.1 Number of unique first names in the database

In the database there are a 100 unique first names stored. This number was determined by taking into consideration that the names in the database were encrypted using deterministic encryption. As a consequence, the number of unique encrypted first names is equivalent to the number of first names in the database.

Listing 1: Java Code to Count Unique First Names

```

1      public int getFirstNameCount() {
2          String query = "SELECT COUNT(DISTINCT enc_firstname) FROM
                           enc_students";
3
4          try (Connection conn = DriverManager.getConnection(url);
5              PreparedStatement pstmt = conn.prepareStatement(query)
6              ;
7              ResultSet rs = pstmt.executeQuery()) {
8              if (rs.next()) {
9                  return rs.getInt(1);
10             }
11
12         } catch (SQLException e) {
13             System.out.println("Error: " + e.getMessage());
14         }
15         return 0;
16     }

```

2.2 Distribution of last names

To plot the distribution of the encrypted last names we determined the number of occurrences of each last name present in the database.

Following, the code to compute the lastname frequency.

Listing 2: Distribution of last names

```

17
18     public HashMap<String, Integer> getLastnameFrequenciesOrdered()
19     {
20         String query = "SELECT enc_lastname FROM enc_students";
21
22         try (Connection conn = DriverManager.getConnection(url);
23             PreparedStatement pstmt = conn.prepareStatement(query)
24             ;
25             ResultSet rs = pstmt.executeQuery()) {
26
27             HashMap<String, Integer> lastnameFrequencies = new
28                 HashMap<>();
29             while (rs.next()) {
30                 String lastname = rs.getString("enc_lastname");
31                 if (lastnameFrequencies.containsKey(lastname)) {
32                     lastnameFrequencies.put(lastname,
33                         lastnameFrequencies.get(lastname) + 1);
34                 } else {
35                     lastnameFrequencies.put(lastname, 1);
36                 }
37             }
38         }
39     }

```

```

32     }
33 }
34
35 return orderHashMapByValue(lastnameFrequencies);
36
37 } catch (SQLException e) {
38     System.out.println("Error: " + e.getMessage());
39 }
40 return null;
41 }

```

Following, the graph with lastnames in chiphertext.




Figure 1: Distribution of encrypted last names.

The graph shows the distribution of the encrypted last names names. The distribution is clearly non-uniform, and an attacker can use background information (like the files given for this assignment) to infer by crafting a lookup table for mapping encrypted values to plaintext values.

Following, the same graph with names in plaintext.




Figure 2: Distribution of encrypted last names.

As it can be seen, lastname Wang (which is the most popular last name) is mapped to the ciphertext

4f09d41d4e6975a14551b684e7ee7739c059bddfd813a8a05a74f88690791445.

2.3 Full names of all students who scored 99 points

There are 100 unique values stored in the database for grades. We also know that the grades can range from 1 to 100. This means that the data set for the grades is a dense dataset, since the complete plaintext space is encrypted. Since the grades are encrypted in an order-preserving manner, the lower numeric value of an encrypted grade corresponds to a lower plaintext value. Based on these information, we sorted the ciphertexts decreasingly, this is how we were able to recover the ciphertext value corresponding to 99.

Following that, we queried all the records where the value of the grade was the previously determined one and we got the students who scored 99 points. The encrypted names could be mapped to the plaintext ones based on the given txt files (again, by leveraging the use of deterministic encryption).

In total there were 149 records with the grade of 99, belonging to 129 students, whose names' are the following: Abdul Yan, Ahmad Cai, Ahmed He, Ahmed Huang, Ahmed Tran, Ali Chen, Ali Martinez, Ana Singh, Ana Zheng, Andrea Mohammad, Andrey Huang, Anna Yang, Antonio Feng, Antonio I, Antonio Xu, Bin Hernandez, Carlos Su, Carlos Yang, Carlos Yu, Carmen Cao, Charles Peng, Daniel Han, David Devi, David Lu, David Zhang, Elena Liu, Elena Wang, Elizabeth Deng, Fatima Khatun, Fatima Wang, Francisco Gao, Francisco Gonzalez, George Kumar, Hong Singh, Hui Cao, Hui Kaur, Hui Liu, Hui Wang, Ibrahim Hassan, Ibrahim Khatun, James Chen, James Zheng, Jean Gao, Joao Ceng, John Li, John Luo, John Singh, John Sun, Jorge Akter, Jorge Chen, Jorge dos Santos, Jose Chen, Jose Li, Jose Tan, Jose Yu, Jose Zhang, Joseph Zhang, Juan Le, Juan Yang, Laura Pan, Lei Hu, Lei Jiang, Lei Wu, Li Li, Li Luo, Ling Xiao, Luis Tan, Manuel Ahmed, Manuel Shi, Maria Li, Maria Wang, Maria Zhang, Marie Gonzalez, Marie Guo, Mario Lin, Mario Yang, Mark Bai, Mark Tan, Martha Tran, Martin Pan, Mary Li, Michael Zhang, Min He, Ming Ahmed, Ming Han, Ming Tang, Mohamed Li, Mohamed Liu, Mohamed Wu, Mohamed Zhang, Mohammed Chen, Mohammed Li, Mohammed Wang, Mohammed Zhang, Muhammad Chen, Muhammad Hu, Muhammad Luo, Nushi Chen, Nushi Li, Nushi Liu, Nushi Wang, Nushi Zhang, Olga Maung, Patrick Zhu, Paul Shi, Peter Luo, Ram Gonzalez, Ram Peng, Richard Liu, Richard Ma, Robert Wu, Rosa Devi, Rosa Wang, Samuel Rodriguez, Samuel Zheng, Sandra Ahmed, Sergey Deng, Sergey Hernandez, Sergey Luo, Siti Khan, Siti Pan, Sri Wu, Wei Chen, Xin Yang, Ying Khan, Ying Wang, Yong Zhu, Yu Deng, Yu Zhu.

Listing 3: Querying for the students who got 99

```

42 public String getSecondUniqueHighestValue() {
43     String query = "SELECT enc_grade FROM ("
44         + "    SELECT DISTINCT enc_grade"
45         + "    FROM enc_students"
46         + "    ORDER BY enc_grade DESC"
47         + ") as eseg LIMIT 1 OFFSET 1;";
48
49     try (Connection conn = DriverManager.getConnection(url);
50         PreparedStatement pstmt = conn.prepareStatement(query)
51         ) {
52
53         // Execute the query
54         ResultSet rs = pstmt.executeQuery();
55
56         // Process the result
57         if (rs.next()) {
58             return rs.getString("enc_grade");
59         }
60     }

```

```

59         } catch (SQLException e) {
60             System.out.println(e.getMessage());
61         }
62         return null;
63     }
64
65     public List<Student> getSecondToTopStudents() {
66         String secondValue = getSecondUniqueHighestValue();
67         String query = "SELECT enc_firstname, enc_lastname FROM
68             enc_students WHERE enc_grade = ?";
69         List<Student> resultList = new ArrayList<>();
70
71         try (Connection conn = DriverManager.getConnection(url);
72             PreparedStatement pstmt = conn.prepareStatement(query)
73             ) {
74
75             pstmt.setString(1, secondValue);
76             ResultSet rs = pstmt.executeQuery();
77
78             ResultSetMetaData metaData = rs.getMetaData();
79             int columnCount = metaData.getColumnCount();
80
81             while (rs.next()) {
82                 Student student = null;
83                 for (int i = 1; i <= columnCount; i++) {
84                     student = new Student(
85                         rs.getString("enc_firstname"),
86                         secondValue,
87                         rs.getString("enc_lastname")
88                     );
89                 }
90                 resultList.add(student);
91             }
92
93         } catch (SQLException e) {
94             System.out.println("Error: " + e.getMessage());
95         }
96
97         return resultList;
98     }

```

■ 2.4 Implementing (α, t) -secure index as mitigation

The main idea behind the (α, t) -secure index is to protect against inference attacks by adding noise to the data in a structured way via clustering the result sets into different partitions. The α parameter specifies the minimum number of keywords that must share a similar access pattern. A higher α means more keywords will be grouped together, increasing privacy but possibly introducing more noise. The t parameter indicates the maximum allowable Hamming distance (i.e., the number of differing bits) between the access patterns of keywords within the same group. A smaller t ensures more similarity between access patterns, improving privacy but potentially leading to more false positives during search results.

We implemented a $(4, 0)$ -secure index by a lookup function, which when searched for by their last names, returns results clusters of 4 keywords. This leads to confusing results when an adversary tries to map the distribution of names similarly to how it

was described previously, but it also introduces multiple false-positive returned by also returning results for other 3 keywords, with which the client executing the query has to parse to identify.

Following, the graph of lastnames' frequencies after implementing a (4,0)-secure index search. As it can be seen, the previous easily mappable distribution is gone, instead clusters of 4 appeared.




Figure 3: Distribution of encrypted last names.

When looking at the plaintext equivalents of the last names from the previous graph, it can also be observed, that only the distribution changed, but so did the few most frequent last names too.




Figure 4: Distribution of the last names in plaintext.

Following, the code

Listing 4: secure index as mitigation

```

97
98 public List<Student> getSecStudentsFromSurname(String surname,
99     int alpha) {
100     List<Student> students = new java.util.ArrayList<>();
101     int i = surnameList.indexOf(surname);
102
103     //compute cluster boundaries
104     int clusterUpperBound = i;
105     for (int j = i; (j % alpha != 0 || j==i) && (j<rows); j++)

```

```

105         {
106             clusterUpperBound = j;
107         }
108         int clusterLowerBound = i;
109         for (int j = i; (j % alpha != 0) && j>0 ; j--) {
110             clusterLowerBound = j-1;
111         }
112
113         //return students in alpha-size cluster
114         if(i != -1) {
115             for (int j = i; j <= clusterUpperBound && (j<rows); j
116                 ++ ) {
117                 for (int k = 0; k < cols; k++) {
118                     if (matrix[j][k]) {
119                         students.add(studentList.get(k));
120                     }
121                 }
122             }
123             i--;
124             for (int j = i; j >= clusterLowerBound; j--) {
125                 for (int k = 0; k < cols; k++) {
126                     if (matrix[j][k]) {
127                         students.add(studentList.get(k));
128                     }
129                 }
130             }
131         }
132         return students;
133     }

```